

THE TRILLION-DOLLAR LEAK : WHY MOST ONLINE SHOPPERS NEVER SAY "YES."

DO YOU KNOW HOW DO THEY OPTIMIZE THE SITE VISIT INTO PURCHASE ..?



Imagine a physical store where 100 people walk in, 66 immediately walk back out after looking at one shelf, and only 33 ever pick up a basket. Of those 33, only a handful make it to the cash register.

In the world of E-commerce, this isn't a disaster—it's the daily reality. Brands spend millions on "Recommendation Engines" designed to nudge users toward that final click. But as user attention spans shrink, the traditional 4-step funnel is dying. Today's shoppers don't want a "journey"; they want the shortest path to "Purchase Confirmed."

When we change the interface to help them, are we actually making it easier—or are we just adding more friction?

A/B Test Report: Recommendation System Interface (EU Region)

What this project focuses on:

This analysis deep-dives into the interface_eu_test, an A/B experiment launched to optimize the shopping experience for new users in the European market.

The Mission:

The goal was clear: Introduce a new recommendation interface and achieve a 10% lift in conversion at every stage of the funnel.

OPTIMIZATION OR FRICTION? A DEEP DIVE INTO THE EU RECOMMENDATION SYSTEM'S FUNNEL PERFORMANCE

A/B Test Report: Recommendation System Interface (EU Region)

1. EXECUTIVE SUMMARY

- **Goal:** Evaluate the effectiveness of a new interface for the recommendation system specifically targeting new EU users.
- **Target KPI:** A minimum **10% increase** in conversion across all funnel stages (product_page, product_cart, and purchase).
- **Result:** The test **did not meet** the expected 10% lift across the entire funnel. While the purchase stage showed a slight absolute improvement of 1.11%, the results were not statistically significant enough to reject the null hypothesis for the overall journey.
- **Recommendation:** Do not roll out the new interface globally. Instead, investigate the high drop-off between the product_page and product_cart and consider optimizing for "one-click" buying behaviours.

2. EXPERIMENTAL DESIGN & METHODOLOGY

- **Test Name:** interface_eu_test.
- **Audience:** New users from the EU region who signed up between Dec 7 and Dec 21, 2020.
- **Groups:**
 - Control (Group A): 4,521 users (current interface).
 - Test (Group B): 4,438 users (new recommendation interface).
- **Analysis Window:** 14 days of data (Dec 7 – Dec 21, 2020)

3. DATA QUALITY & PREPROCESSING

- **Data Integrity:** The final dataset was refined from multiple sources, including a calendar of marketing events, user sign-up logs, and participant assignments.
- **Participant Volume:** After cleaning and ensuring users were correctly attributed to the EU region, the study analyzed **8,959 unique users** performing **51,848 different events**.
- **Splitting Integrity:** Statistical checks confirmed that the population proportions of Groups A and B were split properly, ensuring that any differences in results were likely due to the interface rather than a sampling bias.
- **Sanity Checks Performed:**
 1. Verified no user appeared in both test groups
 2. Confirmed balanced group split
 3. Validated event timestamps and funnel ordering
 4. Confirmed categorical consistency for event types

4. STATISTICAL RESULTS & ANALYSIS

Funnel Stages Evaluated

Login → Product Page → Product Cart → Purchase

Hypothesis:

- H₀: there is not a statistically significant difference in conversion between the samples A and B
- H₁: there is a statistically significant difference in conversion between the samples A and B

The report utilized a series of hypothesis tests (proportions tests) to compare conversion rates at each stage:

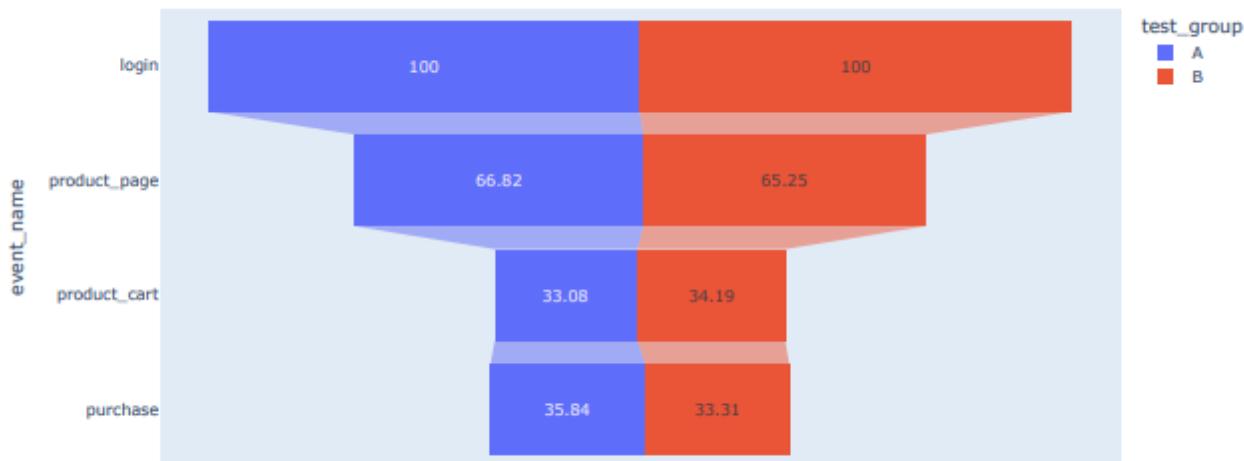
	event_name	p_value	significant	decision
0	login	0.321767	False	Fail to Reject H ₀
1	product_cart	0.147028	False	Fail to Reject H ₀
2	product_page	0.166090	False	Fail to Reject H ₀
3	purchase	0.020365	True	Reject H ₀

Key Insight: While the purchase stage showed significance, it did not reach the target 10% lift. Interestingly, the data suggests some users are bypassing the product_cart entirely, moving directly from the product_page to purchase.

5. DATA QUALITY & FUNNEL INTEGRITY AUDIT

	event_name	n_users	conversion_pct	test_group
0	login	11247	100.00	A
2	product_page	7515	66.82	A
1	product_cart	3720	33.08	A
3	purchase	4031	35.84	A
0	login	10885	100.00	B
2	product_page	7182	65.25	B
1	product_cart	3722	34.19	B
3	purchase	3626	33.31	B

Funnel Analysis- Control Group A vs Test Group B



1. The Logical Funnel Discrepancy

A standard e-commerce funnel requires a linear progression: **Login → Product Page → Product Cart → Purchase**. Mathematically, each subsequent stage must be a subset of the previous stage's unique users.

- **Group B (Test Group):** Exhibits a Healthy Funnel
 - **Conversion Flow:** 100% (Login) → 65.64% (Product Page) → 34.16% (Product Cart) → 33.33% (Purchase).
 - **Observation:** This follows a logical "decay" model. Every step has fewer unique users than the one before it, representing a realistic and clean user journey.

- **Group A (Control Group): Exhibits a Broken Funnel**
- **Conversion Flow:** 100% (Login) → 66.82% (Product Page) → 33.08% (Product Cart) → 35.84% (Purchase).
- **Observation:** The "Purchase" rate is higher than the "Product Cart" rate. It is logically impossible for more unique users to complete a purchase than those who reached the cart stage in a sequential flow. This confirms a structural anomaly in the Control group's data.

2. Statistical Significance vs. Logical Validity

While the Z-test for proportions yielded a "statistically significant" result at the Purchase stage ($p = 0.020$), this result must be interpreted with extreme caution:

- The significance is a by product of the **Control Group's data anomaly**, not necessarily the Test group's performance.
- When the baseline (Group A) is logically inconsistent, any "lift" measured against it is mathematically unreliable.

3. Root Cause Analysis (Technical Hypotheses)

Why is Group A's data defying logic? The discrepancy in the Control Group suggests a failure in the underlying tracking architecture. We have two primary theories for this anomaly:

- **Tracking Failure:** The product_cart event in the legacy interface (Group A) may not be firing consistently across all sessions. This results in an undercount of unique users at the cart stage while still successfully capturing the final purchase event.
- **Legacy Shortcuts:** The old interface may contain a "Buy Now" button or "Direct Checkout" feature that was not accounted for in the initial funnel mapping. This allows users to bypass the cart stage entirely, creating a non-linear journey that the current tracking logic cannot reconcile.

4. Final Strategic Recommendation

Verdict: Inconclusive — Deployment Halted. I recommend against a full rollout of the new interface at this time. To reach a data-driven conclusion, we must:

- **Synchronize Tracking:** Audit and fix the event triggers in the legacy system to ensure both groups are tracked identically.
- **Re-Test:** Re-run the experiment once both funnels are "healthy." This ensures the final decision is based on actual user behavior rather than tracking discrepancies.

6. STRATEGIC RECOMMENDATIONS & NEXT STEPS

Based on the statistical analysis and the funnel integrity audit, I recommend the following:

- **Immediate Action:**

Technical Audit of Control Group (A) The "Purchase > Cart" inversion in Group A is our most significant finding. We must investigate the legacy interface to see if a "Direct Buy" feature exists that is not currently being tracked as a product_cart event.

- **Investigate "Buy Now" as a Global Standard:**

The data from Group A suggests that a significant portion of our users are already finding ways to bypass the cart to complete purchases faster. Instead of forcing users through a traditional recommendation/cart funnel (like in Group B), we should consider making "**Fast-Track Checkout**" a standardized, properly tracked feature for all users.

- **Standardize Event Tracking**

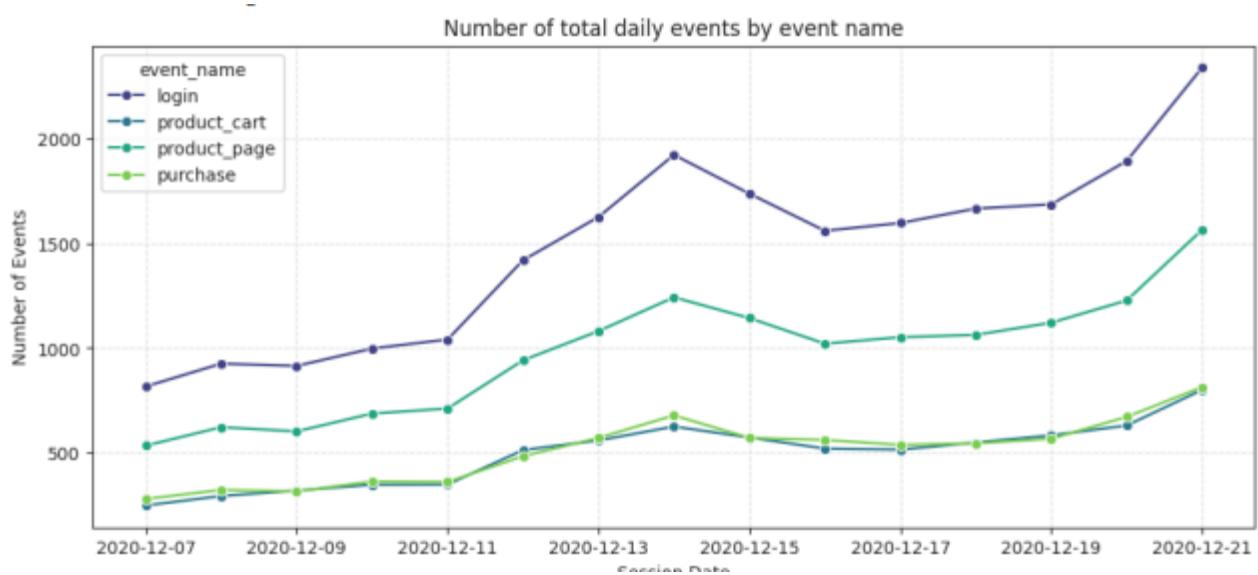
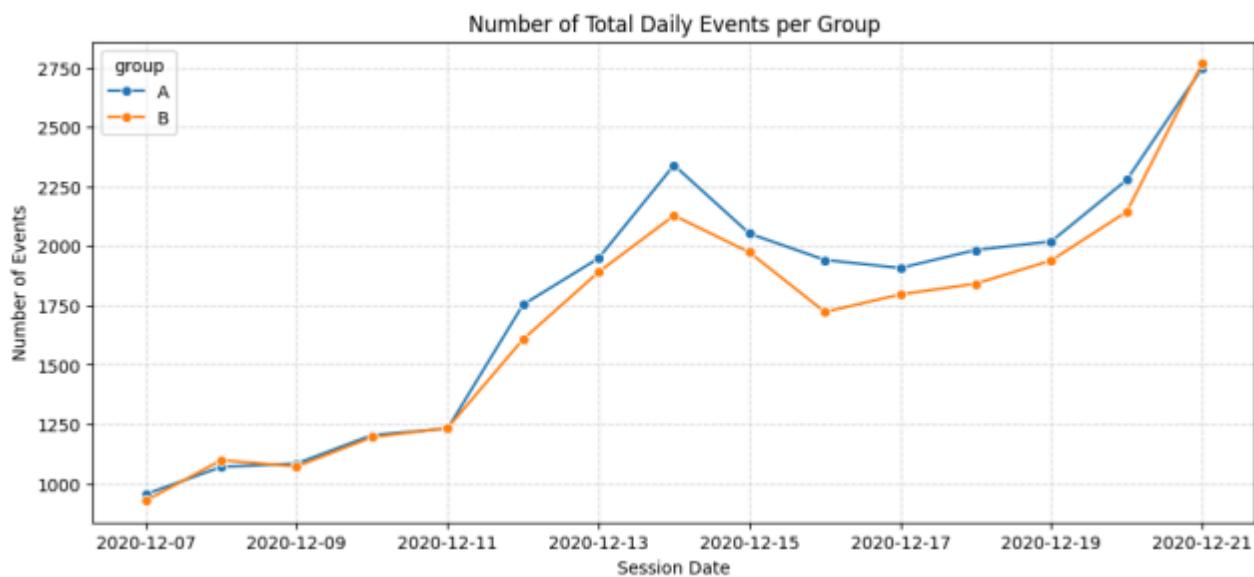
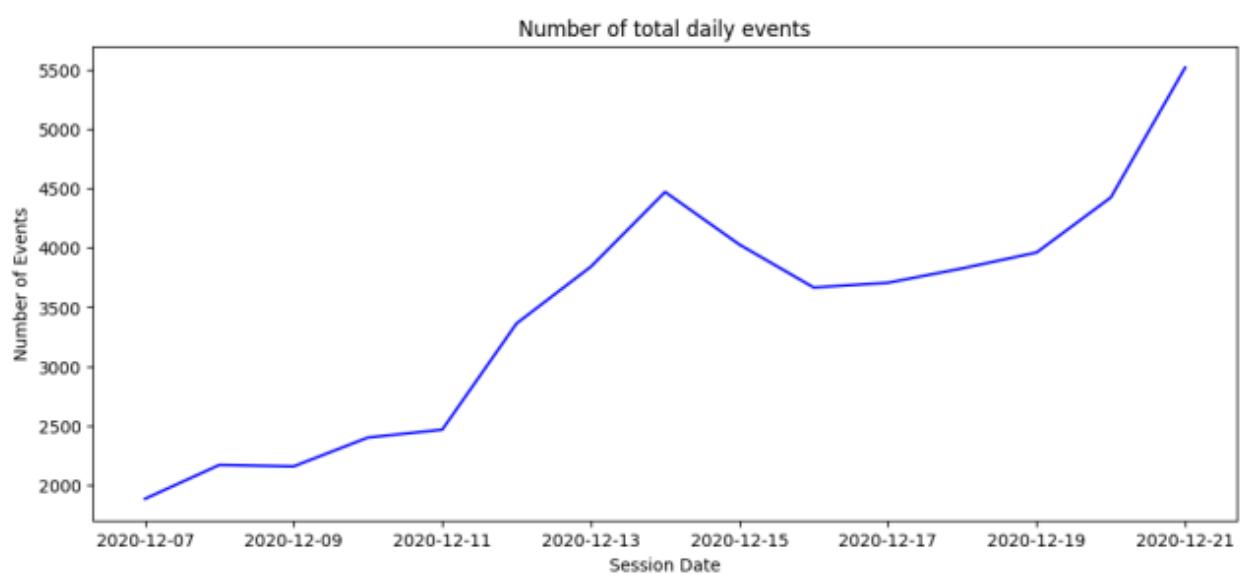
Before re-running this test, ensure that both "Add to Cart" and "Direct Buy" actions are captured as distinct events. This will prevent future funnel discrepancies and allow for a true "apples-to-apples" comparison between new designs and the legacy site.

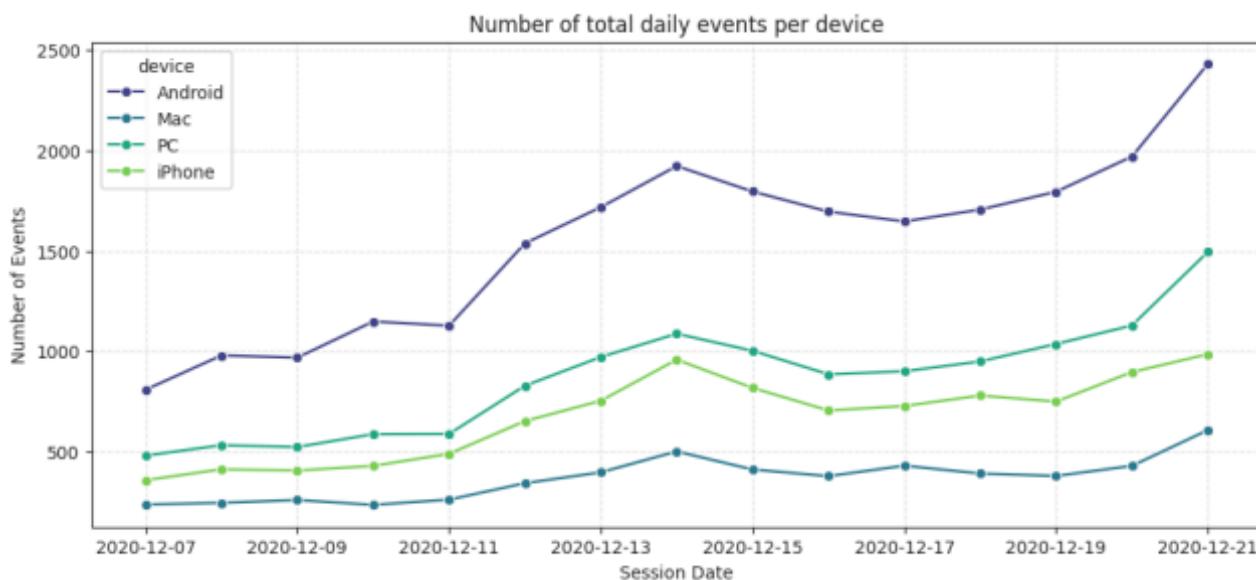
- **Decision:**

Pause Global Rollout We should not deploy the interface_eu_test globally yet. While Group B's funnel is "cleaner," it didn't achieve the 10% lift. More importantly, we cannot make a final decision until we understand why the Control group (A) is outperforming the funnel logic.

7. VISUAL EVIDENCE

How is the number of events distributed by days?





▼ Project description

Project goals

Testing changes:

The introduction of a new interface for an improved recommendation system for EU users.

Test name:

interface_eu_test

Audience:

The audience is new users from the EU who were encouraged to sign up in the online store. Expected number of test participants: 6000

Test duration:

From December 7 to December 21, 2020.

Expected result:

Increase total conversion by at least 10% at each stage of the funnel product_page → product_cart → purchase in test group B.

Data description

- ab_project_marketing_events_us.csv — the calendar of marketing events for 2020
- final_ab_new_users_us.csv — all users who signed up in the online store from December 7 to 21, 2020
- final_ab_events_us.csv — all events of the new users within the period from December 7, 2020 to January 1, 2021
- final_ab_participants_us.csv — table containing test participants

▼ Project Settings

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

▼ Data sources

```
Events = pd.read_csv("https://raw.githubusercontent.com/alona808/ab_testing_online_store/refs/heads/main/datasets/final_ab_"
Marketing_Events = pd.read_csv("https://raw.githubusercontent.com/alona808/ab_testing_online_store/refs/heads/main/datasets/final_
New_Users = pd.read_csv("https://raw.githubusercontent.com/alona808/ab_testing_online_store/refs/heads/main/datasets/final_
Participants = pd.read_csv("https://raw.githubusercontent.com/alona808/ab_testing_online_store/refs/heads/main/datasets/fir_
datasets = [('Events',Events), ('Marketing_Events',Marketing_Events), ('New_Users',New_Users), ('Participants',Participants)
```

▼ Data Overview and Preprocessing

```
for name,dataset in datasets:
    print(f"Dataset_Name : {name}")
    print('      ')
    print(dataset.info())
    print('-----')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 423761 entries, 0 to 423760
Data columns (total 4 columns):
 #   Column       Non-Null Count  Dtype

```

```

0 user_id    423761 non-null object
1 event_dt   423761 non-null object
2 event_name 423761 non-null object
3 details    60314 non-null float64
dtypes: float64(1), object(3)
memory usage: 12.9+ MB
None
-----
Dataset_Name : Marketing_Events

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14 entries, 0 to 13
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   name        14 non-null     object  
 1   regions     14 non-null     object  
 2   start_dt    14 non-null     object  
 3   finish_dt   14 non-null     object  
dtypes: object(4)
memory usage: 580.0+ bytes
None
-----
Dataset_Name : New_Users

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61733 entries, 0 to 61732
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   user_id    61733 non-null     object  
 1   first_date 61733 non-null     object  
 2   region      61733 non-null     object  
 3   device      61733 non-null     object  
dtypes: object(4)
memory usage: 1.9+ MB
None
-----
Dataset_Name : Participants

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18268 entries, 0 to 18267
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   user_id    18268 non-null     object  
 1   group       18268 non-null     object  
 2   ab_test    18268 non-null     object  
dtypes: object(3)
memory usage: 428.3+ KB
None
-----
```

```

for name,dataset in datasets:
    print(f"Dataset_Name : {name}")
    print('    ')
    print(dataset.head())
    print('-----')
```

```

Dataset_Name : Events

      user_id          event_dt event_name  details
0 E1BDDCE0DAFA2679  2020-12-07 20:22:03  purchase   99.99
1 7B6452F081F49504  2020-12-07 09:22:53  purchase    9.99
2 9CD9F34546DF254C  2020-12-07 12:59:29  purchase    4.99
3 96F27A054B191457  2020-12-07 04:02:40  purchase    4.99
4 1FD7660FDF94CA1F  2020-12-07 10:15:09  purchase    4.99
-----
Dataset_Name : Marketing_Events

      name          regions  start_dt \
0  Christmas&New Year Promo      EU, N.America 2020-12-25
1  St. Valentine's Day Giveaway  EU, CIS, APAC, N.America 2020-02-14
2  St. Patric's Day Promo        EU, N.America 2020-03-17
3  Easter Promo                 EU, CIS, APAC, N.America 2020-04-12
4  4th of July Promo            N.America 2020-07-04

      finish_dt
0  2021-01-03
1  2020-02-16
2  2020-03-19
3  2020-04-19
4  2020-07-11
-----
Dataset_Name : New_Users

      user_id  first_date  region  device
0 D72A72121175D8BE  2020-12-07    EU     PC
1 F1C668619DFF6E65  2020-12-07  N.America  Android
2 2E1BF1D4C37EA01F  2020-12-07    EU     PC
```

```
3 50734A22C0C63768 2020-12-07 EU iPhone
4 E1BDDCE0DAFA2679 2020-12-07 N.America iPhone
```

```
-----  
Dataset_Name : Participants
```

```
    user_id group      ab_test
0 D1ABA3E2887B6A73     A recommender_system_test
1 A7A3664BD6242119     A recommender_system_test
2 DABC14FDDFADD29E     A recommender_system_test
3 04988C5DF189632E     A recommender_system_test
4 482F14783456D21B     B recommender_system_test
```

```
-----
```

```
for name,dataset in datasets:
    print(f"Dataset_Name : {name}")
    print('    ')
    print(dataset.describe())
    print('-----')
```

```
Dataset_Name : Events
```

```
    details
count 60314.000000
mean 23.881219
std 72.228884
min 4.990000
25% 4.990000
50% 4.990000
75% 9.990000
max 499.990000
```

```
-----
```

```
Dataset_Name : Marketing_Events
```

```
    name regions start_dt finish_dt
count          14      14      14      14
unique         14       6      14      14
top  Christmas&New Year Promo APAC 2020-12-25 2021-01-03
freq            1       4       1       1
```

```
-----
```

```
Dataset_Name : New_Users
```

```
    user_id first_date region device
count      61733   61733  61733  61733
unique      61733      17      4      4
top  8F04273BB2860229 2020-12-21    EU Android
freq          1      6290  46270  27520
```

```
-----
```

```
Dataset_Name : Participants
```

```
    user_id group      ab_test
count      18268  18268      18268
unique      16666      2          2
top  95401934D6D6D4FC     A interface_eu_test
freq          2      9655      11567
```

```
-----
```

▼ The Marketing Events Dataset

```
Marketing_Events[['start_dt','finish_dt']] = Marketing_Events[['start_dt','finish_dt']].apply(pd.to_datetime)

Marketing_Events=Marketing_Events.sort_values(by='start_dt')

Marketing_Events['n_days']=(Marketing_Events['finish_dt']-Marketing_Events['start_dt']).dt.days

Marketing_Events
```

	name	regions	start_dt	finish_dt	n_days	
6	Chinese New Year Promo	APAC	2020-01-25	2020-02-07	13	
1	St. Valentine's Day Giveaway	EU, CIS, APAC, N.America	2020-02-14	2020-02-16	2	
8	International Women's Day Promo	EU, CIS, APAC	2020-03-08	2020-03-10	2	
2	St. Patrick's Day Promo	EU, N.America	2020-03-17	2020-03-19	2	
3	Easter Promo	EU, CIS, APAC, N.America	2020-04-12	2020-04-19	7	
7	Labor day (May 1st) Ads Campaign	EU, CIS, APAC	2020-05-01	2020-05-03	2	
9	Victory Day CIS (May 9th) Event	CIS	2020-05-09	2020-05-11	2	
11	Dragon Boat Festival Giveaway	APAC	2020-06-25	2020-07-01	6	
4	4th of July Promo	N.America	2020-07-04	2020-07-11	7	
13	Chinese Moon Festival	APAC	2020-10-01	2020-10-07	6	
12	Single's Day Gift Promo	APAC	2020-11-11	2020-11-12	1	
5	Black Friday Ads Campaign	EU, CIS, APAC, N.America	2020-11-26	2020-12-01	5	
0	Christmas&New Year Promo	EU, N.America	2020-12-25	2021-01-03	9	
10	CIS New Year Gift Lottery	CIS	2020-12-30	2021-01-07	8	

Next steps: [Generate code with Marketing_Events](#) [New interactive sheet](#)

Marketing_Events.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 14 entries, 6 to 10
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   name        14 non-null    object  
 1   regions     14 non-null    object  
 2   start_dt    14 non-null    datetime64[ns]
 3   finish_dt   14 non-null    datetime64[ns]
 4   n_days      14 non-null    int64   
dtypes: datetime64[ns](2), int64(1), object(2)
memory usage: 672.0+ bytes
```

▼ Preliminary overview:

- The df_calendar_events dataset contains 14 events over the year 2020.
- The longest event duration is 13 days and belongs to 'Chinese New Year Promo'.
- The shortest event is 1 day. This is 'Single's Day Gift Promo'.
- There are no marketing events in August and September.
- The 'CIS New Year Gift Lottery' event coincides with the 'Christmas&New Year Promo' event. They overlap from December 30 to January 3 inclusive.
- There is no null values.
- The 14 days of A / B testing from December 7 to 21 did not coincide with any marketing event.

▼ Test_Participants_Dataset

Participants[['group', 'ab_test']].unique()

group	ab_test
0	
group	2
ab_test	2

dtype: int64

Participants[['group', 'ab_test']].apply(pd.unique)

group	ab_test
0	recommender_system_test
1	interface_eu_test

```
Participants.duplicated().sum()  
np.int64(0)
```

Preliminary overview:

- The Participants dataset contains 18,268 records
- All users are divided into 2 test groups and 2 types of A / B tests.
- There are no duplicates and Nan values.
- However, there are noticeably fewer unique test participants (16,666) than entries. This means that we have users in the dataset who passed both tests, recommender_system_test and in interface_eu_test.

I think we need to delete users who participate in two tests because we don't know which version of the interface the participants in the second test saw.

```
users_2_tests = (  
    Participants  
    .groupby('user_id')['ab_test']  
    .nunique()  
    .loc[lambda x: x > 1]  
    .reset_index(name='test_count'))  
  
users_2_tests.shape  
(1602, 2)
```

```
users_2_tests.head()
```

	user_id	test_count
0	001064FEAAB631A1	2
1	00341D8401F0F665	2
2	003B6786B4FF5B03	2
3	0082295A41A867B5	2
4	00E68F103C66C1F7	2

Next steps: [Generate code with users_2_tests](#) [New interactive sheet](#)

```
Participants[Participants['user_id']=='0082295A41A867B5']
```

	user_id	group	ab_test
4768	0082295A41A867B5	A	recommender_system_test
14161	0082295A41A867B5	B	interface_eu_test

As we can see, the dataset contains 1602 records of users who participate in both tests. Let's remove them.

```
Participants=Participants[~Participants['user_id'].isin(users_2_tests['user_id'])]  
  
uid = users_2_tests['user_id'].iloc[0]  
Participants[Participants['user_id'] == uid]  
  
Participants.shape[0]
```

```
15064
```

The test_groups dataset has 2 types of A/B test:

recommender_system_test, interface_eu_test. For the purposes of our current project, we have to use data related to interface_eu_test test.

So let's split the dataset to get the necessary data and study it.

```
Participants_EU = Participants[Participants['ab_test']=="interface_eu_test"]  
Participants_EU.shape[0]
```

```
9965
```

```
Participants_by_group = Participants_EU.groupby('group')['user_id'].nunique()
Participants_by_group
```

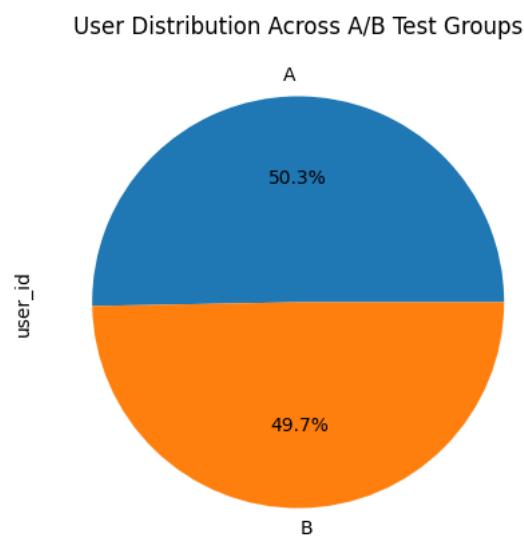
```
user_id
group
A    5012
B    4953
dtype: int64
```

Now let's check if the dataset has users who are in both test groups A and B.

```
EU_2 = Participants_EU.groupby('user_id')['group'].nunique().loc[lambda x: x > 1].reset_index(name='group_count')
EU_2.shape[0]
```

```
0
```

```
Participants_by_group.plot(kind='pie', autopct='%.1f%%', figsize=(5,5), title='User Distribution Across A/B Test Groups')
plt.show()
```



Conclusion:

- After preprocessing, we have a dataset named df_eu_test, which contains the users and groups of data we need, interface_eu_test A / B test.
- The dataset df_eu_test has 9965 unique participants which are almost equally split: * 5012 users - in group A, * 4953 users - in group B
- In our dataset we do not have users that got into both test types or into both test groups.

▼ New_Users_Dataset

```
New_Users.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61733 entries, 0 to 61732
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   user_id     61733 non-null   object 
 1   first_date  61733 non-null   object 
 2   region      61733 non-null   object 
 3   device       61733 non-null   object 
dtypes: object(4)
memory usage: 1.9+ MB
```

Count unique values in the columns:

```
New_Users.nunique()
```

```

0
user_id    61733
first_date   17
region      4
device       4

dtype: int64

```

```
New_Users[['region', 'device']].agg(pd.unique)
```

	region	device	grid icon
0	EU	PC	
1	N.America	Android	
2	APAC	iPhone	
3	CIS	Mac	

```

New_Users.duplicated().sum()
np.int64(0)

```

Preliminary overview:

- The new_users dataset contains 61733 unique entries.
- New users signed up in the online store from December 7 to 23, 2020.
- New users have come from 4 regions: EU, N.America, APAC, CIS.
- They utilize 4 types of devices: PC, Android, iPhone, Mac.
- There are no duplicates and Nan values.

To test the new interface of the online store, we have 2 mandatory testing conditions:

- Testing lasted 14 days, from December 7 to December 21. So these days are included in the scope of the project;
- We have to research new EU users.

Firstly, let's look at the days of December 22nd and 23rd as outside the scope of our task and remove them.

```

New_Users= New_Users[New_Users['first_date']<"2020-12-22"]

New_Users['first_date'].max()
'2020-12-21'

```

Secondly, we have to extract data on new EU users from the dataset. Moreover, they must be our test participants.

```

New_Users_EU = New_Users[(New_Users['user_id'].isin (Participants_EU['user_id'])) & (New_Users['region']=='EU')]
New_Users_EU.shape[0]

8963

```

The dataset of new EU users who signed up in the online store from December 7 to 21, 2020 contains fewer unique users than the dataset containing test participants:

- dataset df_new_eu_users: 8963 users
- dataset df_eu_test: 9965 users

I can assume that some users took part in A / B testing without prior registration in the online store, for example, using the 'Buy Now' button, etc. We can check this version later.

So let's continue with our new EU users.

```

Daily_New_Users = (New_Users_EU.groupby('first_date')['user_id'].nunique())
print(Daily_New_Users)

Daily_New_Users.plot(kind='line', figsize=(10,5), title='Daily Number of Newly Signed Users', color='blue', xlabel='Session Date')
plt.show()

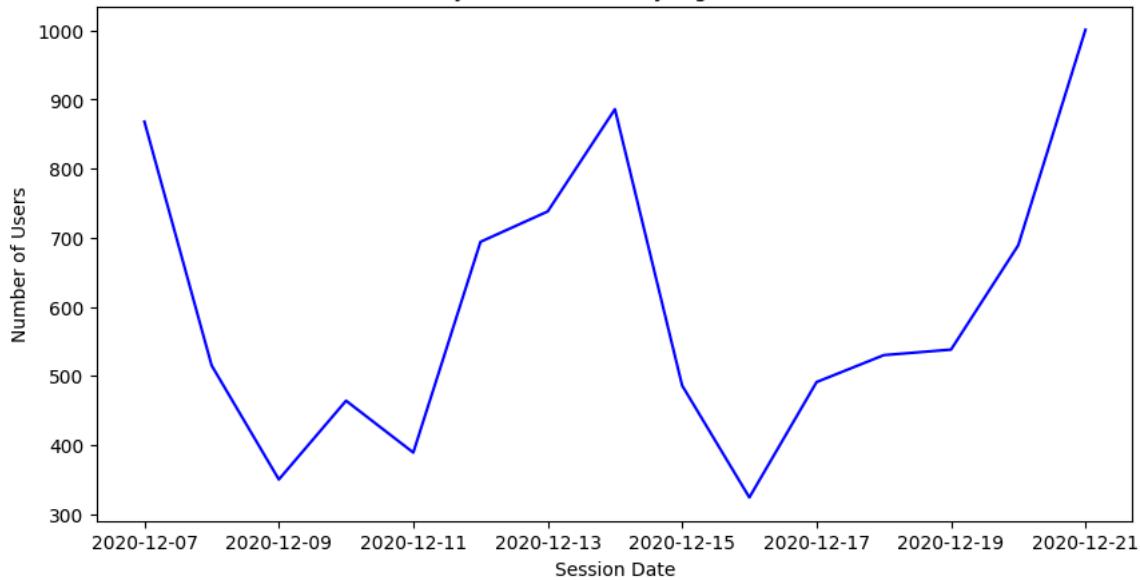
```

```

first_date
2020-12-07    868
2020-12-08    515
2020-12-09    350
2020-12-10    464
2020-12-11    389
2020-12-12    694
2020-12-13    738
2020-12-14    886
2020-12-15    486
2020-12-16    324
2020-12-17    491
2020-12-18    530
2020-12-19    538
2020-12-20    689
2020-12-21   1001
Name: user_id, dtype: int64

```

Daily Number of Newly Signed Users



- However, in fact, EU users make up 82% of our dataset.

```

New_Users_by_Device = (New_Users_EU.groupby('device')['user_id'].nunique())
print(New_Users_by_Device)

```

```

New_Users_by_Device.plot(kind='pie', autopct='%1.1f%%', figsize=(5,5), title='User Distribution Across Devices')
plt.show()

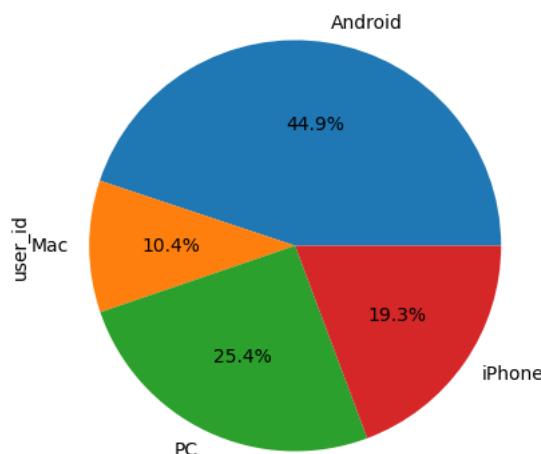
```

```

device
Android    4022
Mac         930
PC          2281
iPhone      1730
Name: user_id, dtype: int64

```

User Distribution Across Devices



Conclusion:

- After preprocessing we have the df_new_eu_users dataset 8963 unique entries.
- All new users in the dataset are from the EU who signed up in the online store from December 7 to 21, 2020.
- New EU user registrations for 14 days of testing varied greatly:
 - The maximum number of new users came on December 21st - 1001 new users,
 - The minimum - on December 16th - 324 new users,
 - The number of new user registrations depends on the day of week.
- They utilize 4 types of devices: PC, Android, iPhone, Mac:
 - Almost half of the new users are Android users, 44.9%.
 - The smallest number of users came using the Mac, 10.4%.
- There are no null values and duplicates in the dataset.
- According to the assignment, the estimated number of test participants must be at least 6000. We have 8963 users, which meets the conditions of testing.

Events dataset

```
print(Events.shape)
Events.head()

(423761, 4)
   user_id      event_dt  event_name  details
0  E1BDDCE0DAFA2679  2020-12-07 20:22:03  purchase    99.99
1  7B6452F081F49504  2020-12-07 09:22:53  purchase     9.99
2  9CD9F34546DF254C  2020-12-07 12:59:29  purchase     4.99
3  96F27A054B191457  2020-12-07 04:02:40  purchase     4.99
4  1FD7660FDF94CA1F  2020-12-07 10:15:09  purchase     4.99
```

Count unique values in the columns:

```
Events.nunique()
```

	0
user_id	58703
event_dt	257138
event_name	4
details	4

```
dtype: int64
```

```
Events[['event_name','details']].agg(pd.unique)
```

	0
event_name	[purchase, product_cart, product_page, login]
details	[99.99, 9.99, 4.99, 499.99, nan]

```
dtype: object
```

Let's explore Nan values in the details column:

```
details_null = Events['details'].isna().sum()
details_null_pct = (Events['details'].isna().mean() * 100).round(2)

details_null, details_null_pct

(np.int64(363447), np.float64(85.77))
```

As we can see, our event dataset has 85.77% NaN values in the details column.

I think this happened because not all of our users did reach the purchase event. We can check this version)

```
Events[(Events['event_name']=='purchase') & (Events['details'].isna())].shape[0]
```

```
0
```

There is no 'purchase' event with NaN value in the 'details' column. But previous steps of funnel have Nan's.

I consider that we can fill Nan's with 0.

```
Events.details.fillna(0,inplace=True)  
Events.isna().sum()
```

```
0  
user_id    0  
event_dt   0  
event_name  0  
details     0  
  
dtype: int64
```

```
Events.duplicated ().sum()  
np.int64(0)
```

```
Events  
Events.event_dt.describe()
```

```
event_dt  
count      423761  
unique     257138  
top        2020-12-23 02:37:24  
freq       10  
  
dtype: object
```

```
Events['event_dt']=pd.to_datetime(Events['event_dt'])  
Events=Events.sort_values(by='event_dt')  
Events['event_dt'].describe()
```

```
event_dt  
count      423761  
mean      2020-12-18 10:10:17.282395136  
min       2020-12-07 00:00:33  
25%      2020-12-14 03:05:18  
50%      2020-12-18 17:40:52  
75%      2020-12-22 13:51:17  
max       2020-12-30 23:36:33  
  
dtype: object
```

Preliminary overview:

- The events dataset contains 423761 records from December 7 to December 30, 2020.
- The funnel consists 4 events: login, product_page, product_cart, and purchase. However, in accordance with our task, we are interested in a funnel of these 3 steps: product_page → product_card → purchase.
- There are no duplicates.
- Nan values of the variable details were filled with 0.

To prepare the dataset for further exploration, we need to perform the same data transformation as for the new user dataset, namely:

- First, get the data for our testing from December 7 to December 21;
- Second, we need to retrieve the events that were executed by our test participants.*

Let's create the df_test_events for our goals and explore it:

```
Events_EU = Events[(Events['user_id'].isin(Participants_EU['user_id'])) & (Events['event_dt'] < "2020-12-22")]
Events_EU.shape[0]
```

```
51848
```

Count unique values in the columns:

```
Events_EU.nunique()
```

	0
user_id	8959
event_dt	33338
event_name	4
details	5

```
dtype: int64
```

```
print (Events_EU.sample(5))
```

```
print('  ')
```

```
print (Events_EU.event_name.value_counts())
```

```
print('  ')
```

```
Events_EU.duplicated ().sum()
```

	user_id	event_dt	event_name	details
196620	CE8B1BCBCA042768	2020-12-20 19:52:29	product_page	0.0
339043	233653C0930FD807	2020-12-19 17:02:47	login	0.0
265953	80BE72C98A257C69	2020-12-11 10:20:44	login	0.0
285678	7B125EC3570700E	2020-12-14 16:52:24	login	0.0
293795	8155910F11B56B21	2020-12-14 10:27:23	login	0.0

	event_name	count
login	22132	
product_page	14617	
purchase	7657	
product_cart	7442	

```
Name: count, dtype: int64
```

```
np.int64(0)
```

First look at the number of events per user:

```
Events_EU_per_user = Events_EU.groupby('user_id')['event_dt'].count()
Events_EU_per_user.describe()
```

	event_dt
count	8959.000000
mean	5.787253
std	3.727047
min	1.000000
25%	3.000000
50%	6.000000
75%	8.000000
max	24.000000

```
dtype: float64
```

```
Events_EU_per_user.mode()
```

	event_dt
0	6

```
dtype: int64
```

Conclusion

- The df_test_events dataset contains 51848 events.
- The dataset includes data that was collected from December 7 to December 21, 2020.
- All events presented in the dataset were performed by test takers.
- Each participant completed at least 1 event. Maximum events performed by user is 24.
- An average number of events per user is 6.
- Most often the number of events per user is 6, which coincides with the average number of events.

✓ Merging Datasets

After preprocessing our datasets, we can combine them for deeper exploratory analysis. But first, a small reminder of which ones we will use and what they look like :)

Participants_EU.head()

	user_id	group	ab_test	grid icon
6701	D4E530F6595A05A3	A	interface_eu_test	
6703	6BCB0F33D3BAB8C2	A	interface_eu_test	
6707	2D2E7AA539CF348F	B	interface_eu_test	
6708	4BA448BCE1343C6F	A	interface_eu_test	
6709	76B6CDF2A8B1DBFC	B	interface_eu_test	

Next steps: [Generate code with Participants_EU](#) [New interactive sheet](#)

New_Users_EU.head()

	user_id	first_date	region	device	grid icon
2	2E1BF1D4C37EA01F	2020-12-07	EU	PC	
3	50734A22C0C63768	2020-12-07	EU	iPhone	
14	5BE017E9C8CC42F8	2020-12-07	EU	Android	
18	96F27A054B191457	2020-12-07	EU	iPhone	
21	E6AF85675078215D	2020-12-07	EU	Android	

Events_EU.head()

	user_id	event_dt	event_name	details	grid icon
245271	B13A53A1EB2038EE	2020-12-07 00:02:48	login	0.0	
245598	4C4BA430AAA820F8	2020-12-07 00:03:18	login	0.0	
123266	4C4BA430AAA820F8	2020-12-07 00:03:19	product_page	0.0	
242531	4B7C59A60FE1DA69	2020-12-07 00:03:51	login	0.0	
60719	4B7C59A60FE1DA69	2020-12-07 00:03:54	product_cart	0.0	

Next steps: [Generate code with Events_EU](#) [New interactive sheet](#)

```
print('Participants_EU :', Participants_EU.shape[0], ' New_users_EU :', New_Users_EU.shape[0], ' Events_EU :', Events_EU.shape[0])
Participants_EU : 9965  New_users_EU : 8963  Events_EU : 51848
```

Combining the datasets:

```
df_ab_testing = Events_EU.merge(Participants_EU[['user_id', 'group']], on='user_id', how='left').merge(New_Users_EU[['user_id', 'region', 'device']], on='user_id', how='left')
```

	user_id	event_dt	event_name	details	group	first_date	device	grid icon
0	B13A53A1EB2038EE	2020-12-07 00:02:48	login	0.0	A	2020-12-07	PC	
1	4C4BA430AAA820F8	2020-12-07 00:03:18	login	0.0	A	2020-12-07	PC	
2	4C4BA430AAA820F8	2020-12-07 00:03:19	product_page	0.0	A	2020-12-07	PC	
3	4B7C59A60FE1DA69	2020-12-07 00:03:51	login	0.0	B	2020-12-07	Android	
4	4B7C59A60FE1DA69	2020-12-07 00:03:54	product_cart	0.0	B	2020-12-07	Android	

Next steps: [Generate code with df_ab_testing](#) [New interactive sheet](#)

```
df_ab_testing.isnull().sum()
```

	0
user_id	0
event_dt	0
event_name	0
details	0
group	0
first_date	0
device	0

dtype: int64

```
df_ab_testing.duplicated().sum()
```

```
np.int64(0)
```

▼ Exploratory Data Analysis (EDA)

▼ Study conversion at different funnel stages.

The main goal of our test is to make sure that the total conversion [1] on each stage in Test Group B is better than the total conversion at the same stage in Control Group A.

[1] 'Total conversion is the ratio of users at a particular stage to the number of users at the first stage.'

Let's plot the funnel and compare the groups.

1. Calculate the number of users for each stage in group A:
2. Calculating total conversion for each funnel stage in group A:

```
funnel_order=['login','product_page','product_cart','purchase']

df_funnel_A = (
    df_ab_testing[df_ab_testing['group'] == 'A']
    .groupby('event_name')['user_id']
    .count()
    .reset_index(name='n_users')
)

# enforce logical order
df_funnel_A['event_name'] = pd.Categorical(
    df_funnel_A['event_name'],
    categories=funnel_order,
    ordered=True)

df_funnel_A = df_funnel_A.sort_values('event_name')

first_stage_A = df_funnel_A['n_users'].iloc[0]
print(f'First Stage Users : {first_stage_A}')

print('')

df_funnel_A['conversion_pct'] = (df_funnel_A['n_users'] / first_stage_A * 100).round(2)
df_funnel_A['test_group'] = 'A'
print(df_funnel_A)
```

```
First Stage Users : 11247
```

	event_name	n_users	conversion_pct	test_group
0	login	11247	100.00	A
2	product_page	7515	66.82	A
1	product_cart	3720	33.08	A
3	purchase	4031	35.84	A

The same steps for the group B:

```
df_funnel_B = (
    df_ab_testing[df_ab_testing['group'] == 'B']
    .groupby('event_name')['user_id']
    .count()
    .reset_index(name='n_users'))

df_funnel_B['event_name'] = pd.Categorical(
    df_funnel_B['event_name'],
    categories=funnel_order,
    ordered=True)

df_funnel_B = df_funnel_B.sort_values('event_name')

first_stage_B = df_funnel_B['n_users'].iloc[0]
print(f'First Stage Users : {first_stage_B}')

print('      ')

df_funnel_B['conversion_pct'] = (df_funnel_B['n_users'] / first_stage_B * 100).round(2)
df_funnel_B['test_group'] = 'B'
print(df_funnel_B)
```

```
First Stage Users : 10885
```

	event_name	n_users	conversion_pct	test_group
0	login	10885	100.00	B
2	product_page	7102	65.25	B
1	product_cart	3722	34.19	B
3	purchase	3626	33.31	B

Concatenating test groups and plotting funnel chart:

```
df_funnel = pd.concat([df_funnel_A, df_funnel_B], axis=0)
print(df_funnel)

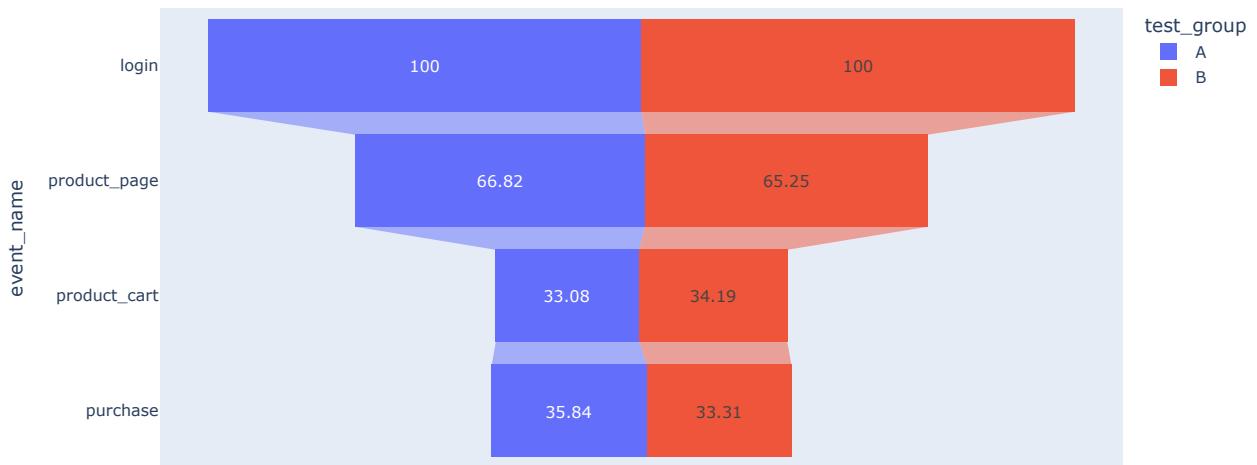
print('      ')

import plotly.express as px

Funnel = px.funnel(df_funnel, x='conversion_pct', y='event_name', color='test_group', title='Funnel Analysis- Control Group A')
Funnel.show()
```

	event_name	n_users	conversion_pct	test_group
0	login	11247	100.00	A
2	product_page	7515	66.82	A
1	product_cart	3720	33.08	A
3	purchase	4031	35.84	A
0	login	10885	100.00	B
2	product_page	7102	65.25	B
1	product_cart	3722	34.19	B
3	purchase	3626	33.31	B

Funnel Analysis- Control Group A vs Test Group B



- According to the test, it was expected that 'At each stage of the funnel product_page → product_cart → purchase, there will be at least a 10% increase.'
- However, the funnel chart shows that the conversion in the control group A is even slightly higher than in the test group B in the first two stages.
- Only the Product_Cart stage showed 1.11% better in test group B than in group A.
- Moreover, at purchase stage in group A, the total conversion is higher than at the previous stage. I assume this was due to the 'Buy Now' button, which allowed users to buy by skipping some steps.

Is the number of events per user distributed equally in the samples?

```

summary = df_ab_testing.groupby('group').agg(
    n_users=('user_id', 'nunique'),
    n_events=('event_dt', 'count')
).reset_index()

summary['events_per_user'] = (summary['n_events'] / summary['n_users']).round(2)

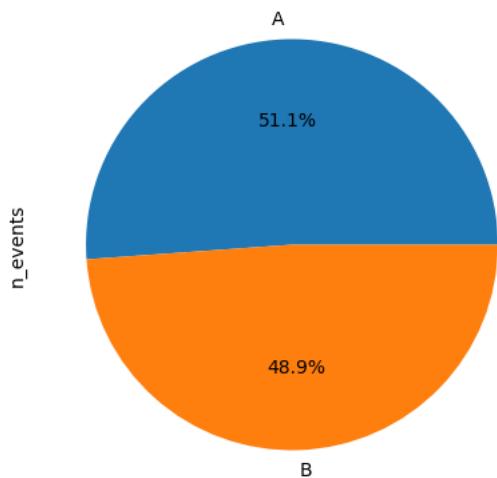
print(summary)

summary['n_events'].plot(kind='pie', labels=summary['group'], autopct='%1.1f%%', figsize=(5,5), title='Events Distribution Across Groups')
plt.show()

```

	group	n_users	n_events	events_per_user
0	A	4521	26513	5.86
1	B	4438	25335	5.71

Events Distribution Across Groups



As we can see, the share of users and events is close to 50/50.

Thus, we can claim that the number of events per user distributed equally in the samples

Are there users who enter both samples?

```
users_in_multiple_groups = (
    df_ab_testing
    .groupby('user_id')['group']
    .nunique()
)

users_in_multiple_groups[users_in_multiple_groups > 1].shape[0]
```

0

There is no users who enter both samples.

How is the number of events distributed by days?

```
df_ab_testing['event_date']=df_ab_testing['event_dt'].dt.date
df_ab_testing.head()
```

	user_id	event_dt	event_name	details	group	first_date	device	event_date	grid icon
0	B13A53A1EB2038EE	2020-12-07 00:02:48		login	0.0	A	2020-12-07	PC	2020-12-07
1	4C4BA430AAA820F8	2020-12-07 00:03:18		login	0.0	A	2020-12-07	PC	2020-12-07
2	4C4BA430AAA820F8	2020-12-07 00:03:19	product_page		0.0	A	2020-12-07	PC	2020-12-07
3	4B7C59A60FE1DA69	2020-12-07 00:03:51		login	0.0	B	2020-12-07	Android	2020-12-07
4	4B7C59A60FE1DA69	2020-12-07 00:03:54	product_cart		0.0	B	2020-12-07	Android	2020-12-07

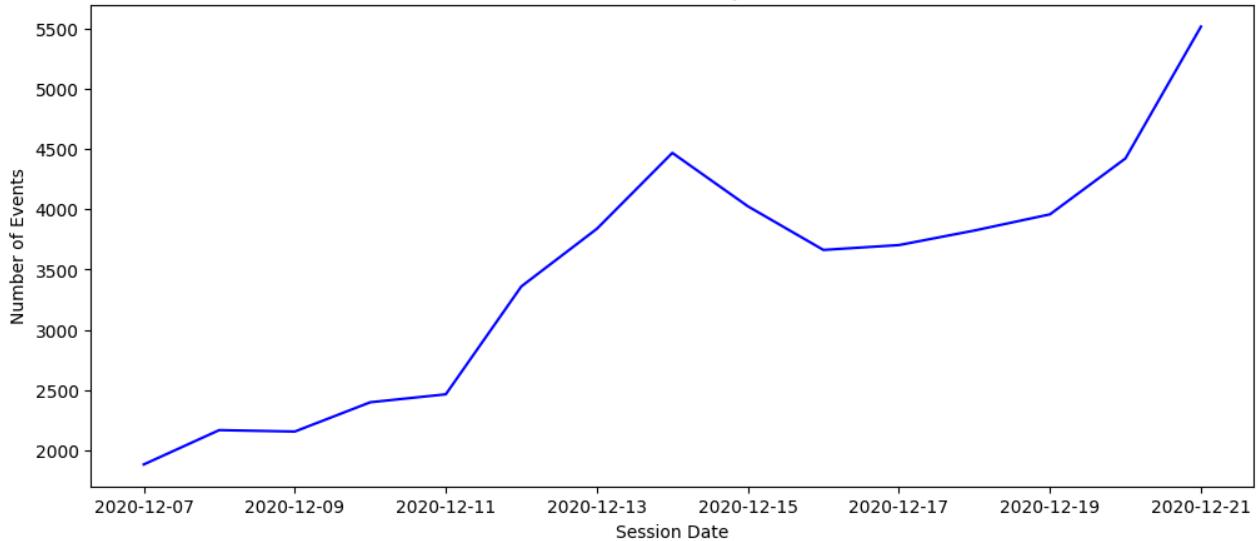
Next steps: [Generate code with df_ab_testing](#) [New interactive sheet](#)

Number of total daily events:

```
Events_per_day = df_ab_testing.groupby('event_date')['user_id'].count()

Events_per_day.plot(kind='line', figsize=(12,5), title='Number of total daily events', color='blue', xlabel='Session Date', ylab
plt.show()
```

Number of total daily events



Number of total daily events per test groups:

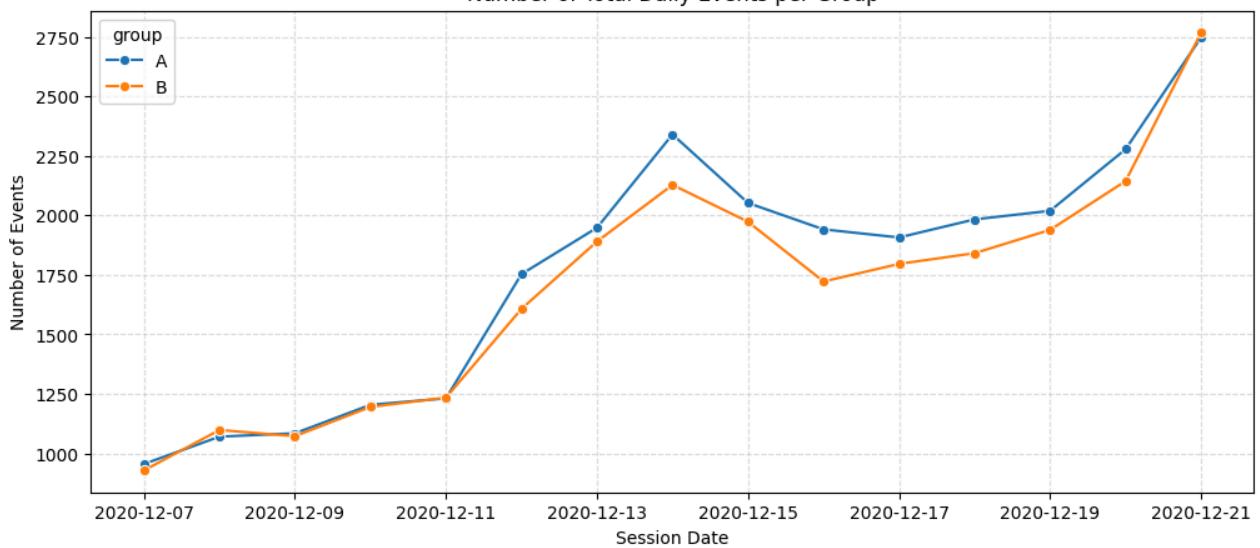
```
Events_per_group = df_ab_testing.groupby(['event_date','group'])['event_dt'].count().reset_index(name = 'n_events')
print(Events_per_group.head())

plt.figure(figsize=(12,5))
sns.lineplot(data=Events_per_group,x='event_date',y='n_events',hue='group',marker='o')
plt.title('Number of Total Daily Events per Group')
plt.xlabel('Session Date')
plt.ylabel('Number of Events')

sns.set_style("ticks")
plt.grid(True, linestyle='--', alpha=0.4)
plt.show()
```

	event_date	group	n_events
0	2020-12-07	A	955
1	2020-12-07	B	928
2	2020-12-08	A	1070
3	2020-12-08	B	1098
4	2020-12-09	A	1084

Number of Total Daily Events per Group



Number of total daily events by event name:

```
Events_by_eventname = df_ab_testing.groupby(['event_date','event_name'])['event_dt'].count().reset_index(name = 'n_events')
print(Events_by_eventname.head())
```

```

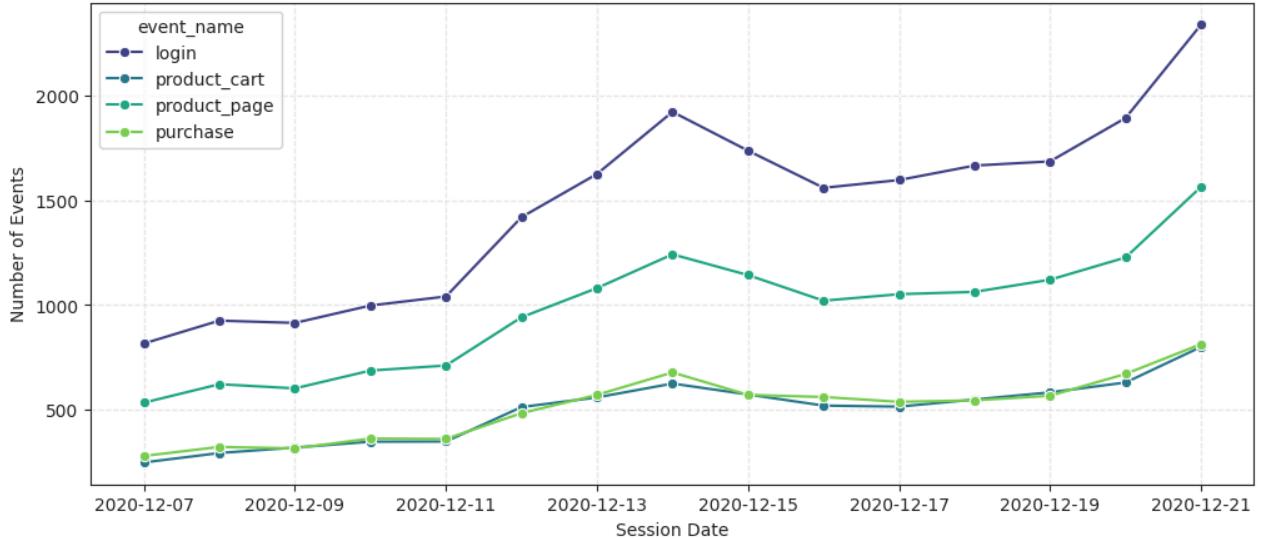
plt.figure(figsize=(12,5))
sns.lineplot(data=Events_by_eventname,x='event_date',y='n_events',hue='event_name',marker='o',palette='viridis')
plt.title('Number of total daily events by event name')
plt.xlabel('Session Date')
plt.ylabel('Number of Events')

sns.set_style("ticks")
plt.grid(True, linestyle='--', alpha=0.4)
plt.show()

```

	event_date	event_name	n_events
0	2020-12-07	login	817
1	2020-12-07	product_cart	250
2	2020-12-07	product_page	535
3	2020-12-07	purchase	281
4	2020-12-08	login	926

Number of total daily events by event name



Number of total daily events by device:

```

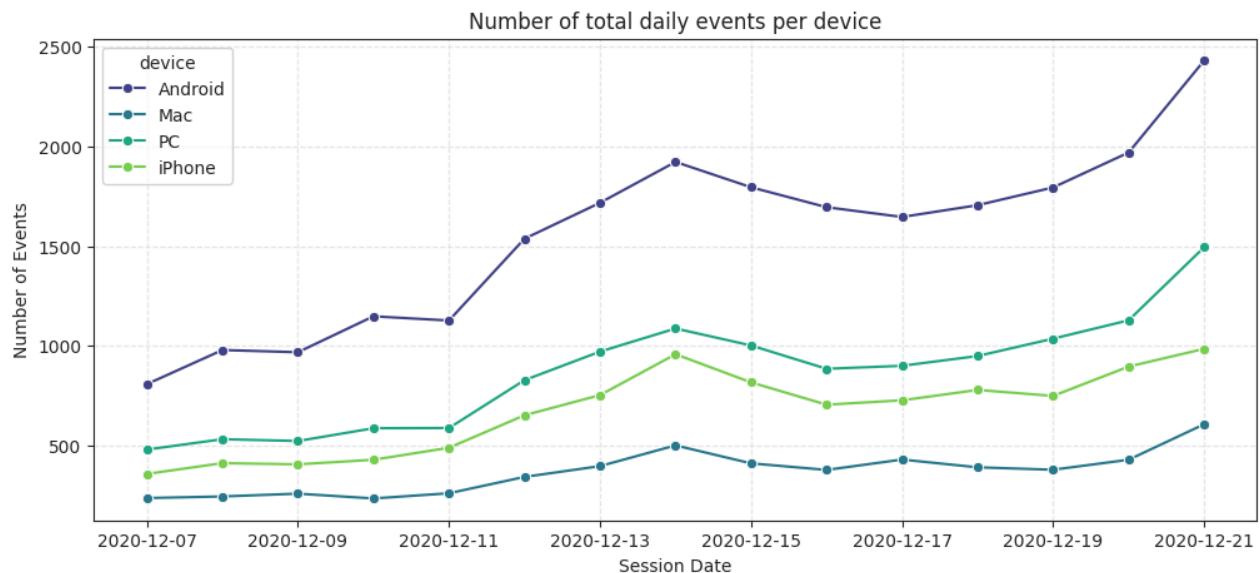
Events_by_device = df_ab_testing.groupby(['event_date','device'])['event_dt'].count().reset_index(name = 'n_events')
print(Events_by_device.head())

plt.figure(figsize=(12,5))
sns.lineplot(data=Events_by_device,x='event_date',y='n_events',hue='device',marker='o',palette='viridis')
plt.title('Number of total daily events per device')
plt.xlabel('Session Date')
plt.ylabel('Number of Events')

sns.set_style("ticks")
plt.grid(True, linestyle='--', alpha=0.4)
plt.show()

```

	event_date	device	n_events
0	2020-12-07	Android	808
1	2020-12-07	Mac	237
2	2020-12-07	PC	480
3	2020-12-07	iPhone	358
4	2020-12-08	Android	979



- The total number of events during the test increased significantly, 3 times.
- This feature is typical for both Groups A and B, for each event name and for number of events per device.
- Most of the events were performed by Android users, with the smallest user group being Mac fans.
- It is clear from the graphs that we can say that there were spikes of events on a certain weekday. In our case, it was on Thursdays.
- To funnel stages, product_cart and purchase have almost the same number of events during the test weeks. I can assume that the abandoned cart percentage is very low, or the 'Buy Now' button works better than the whole funnel.

Think of the possible details in the data that you have to take into account before starting the A/B test?

- The datasets had users who were assigned to multiple groups and who performed both tests at the same time. We have processed our data to avoid exposure to these factors.
- We also checked if the users are equally split.
- But there is an external factor that influences the test result. I consider that we have to take into account that the A / B test was carried out on the eve of Christmas. And even if the specific testing time did not last at the same time as the marketing activities, we cannot trust the results 100% due to the well-known increase in consumer activity before the Christmas holidays.

Conclusion:

- The test group B did not achieve expected result of increasing total conversion at least 10% on each stage of the funnel product_page → product_cart → purchase.
- The test group B showed better result on the purchase stage.
- There is no users who enter both samples.
- The number of events per user distributed equally in the samples.

Evaluate the A/B test results

What can you tell about the A/A test results?

```
df_users_pivot = df_ab_testing.pivot_table(index='event_name', columns='group', values='user_id', aggfunc='nunique').reset_index()
df_users_pivot
```

group	event_name	A	B	
0	login	4520	4438	edit
1	product_cart	1479	1516	
2	product_page	3030	2913	
3	purchase	1612	1479	

Next steps: [Generate code with df_users_pivot](#) [New interactive sheet](#)

```
df_users_pivot['diff_n_users'] = df_users_pivot['A'] - df_users_pivot['B']
df_users_pivot
```

group	event_name	A	B	diff_n_users	
0	login	4520	4438	82	edit
1	product_cart	1479	1516	-37	
2	product_page	3030	2913	117	
3	purchase	1612	1479	133	

Next steps: [Generate code with df_users_pivot](#) [New interactive sheet](#)

```
df_users_crosstab = pd.crosstab(index= df_ab_testing['event_name'], columns=df_ab_testing['group'],values=df_ab_testing['us'],
aggfunc='nunique',normalize=True,margins=True,margins_name='Total').reset_index()
```

df_users_crosstab

group	event_name	A	B	Total	
0	login	0.215371	0.211464	0.426836	edit
1	product_cart	0.070472	0.072235	0.142707	
2	product_page	0.144375	0.138800	0.283175	
3	purchase	0.076809	0.070472	0.147282	
4	Total	0.504632	0.495368	1.000000	

Next steps: [Generate code with df_users_crosstab](#) [New interactive sheet](#)

```
df_users_crosstab['diff_share_stage']=((df_users_crosstab['A']-df_users_crosstab['B'])*100).round(2)
df_users_crosstab
```

group	event_name	A	B	Total	diff_share_stage	
0	login	0.215371	0.211464	0.426836	0.39	edit
1	product_cart	0.070472	0.072235	0.142707	-0.18	
2	product_page	0.144375	0.138800	0.283175	0.56	
3	purchase	0.076809	0.070472	0.147282	0.63	
4	Total	0.504632	0.495368	1.000000	0.93	

Next steps: [Generate code with df_users_crosstab](#) [New interactive sheet](#)

We can argue that all the criteria we have for a successful A/A test:

- As can be seen from the aggregated tables, there is no significant difference between the samples. On Each stage the difference is less than 1%.
- For all groups, data on the same event is recorded.
- Users remain within their groups until the end of the test.
- Users who may see different versions of the page during the test were removed from the groups.

Use the z-criterion to check the statistical difference between the proportions

H0: there is not a statistically significant difference in conversion between the samples A and B

H1: there is a statistically significant difference in conversion between the samples A and B

```
event = df_users_pivot['event_name'].unique()    ## Preparing event data
alpha = 0.05
results=[]
from statsmodels.stats.proportion import proportions_ztest

for eve in event:      ## Loop for all event Z-test
    a_sucess = df_users_pivot.loc[df_users_pivot['event_name']==eve,'A'].iloc[0]
    b_sucess = df_users_pivot.loc[df_users_pivot['event_name']==eve,'B'].iloc[0]

    a_total=df_ab_testing[df_ab_testing['group']=='A']['user_id'].nunique()
    b_total=df_ab_testing[df_ab_testing['group']=='B']['user_id'].nunique()

    z_stat,p_value=proportions_ztest(count=[a_sucess,b_sucess],nobs=[a_total,b_total])

    results.append({'event_name':eve,'p_value':round(p_value,6),'significant':p_value < alpha,'decision': 'Reject H0' if p_value < alpha else 'Fail to Reject H0'})
```

```
ab_test_results = pd.DataFrame(results)
ab_test_results
```

	event_name	p_value	significant	decision
0	login	0.321767	False	Fail to Reject H ₀
1	product_cart	0.147028	False	Fail to Reject H ₀
2	product_page	0.166090	False	Fail to Reject H ₀
3	purchase	0.020365	True	Reject H ₀

Next steps: [Generate code with ab_test_results](#) [New interactive sheet](#)

Conclusion:

- As a result of checking the hypothesis, we cannot reject H₀ for all the events.
- Thus, we can accept the equality of the proportions of the population of A and B groups. It means that the groups were split properly.

General conclusion

Datasets:

- The logs were collected in several tables for the period from December 7 to December 30, 2020.
- The two types of tests were run simultaneously and overlapped using the same test participants.

New interface for EU users test:

- The test was carried out within 14 days, from December 7 to December 21, 2020.
- After preprocessing, we got a dataset that contains 51848 different events performed by 8959 unique users.
- Users are split equally into two groups: - A control group - 4521 users - B test group - 4438 users
- Control group A gets the old site version and one test group B gets the new ones:

The funnel:

- Consists of 4 events: login → product_page → product_cart → purchase
- The number of events per user distributed equally in the samples.

Performing tests:

- The A/A test shows that the control groups were split properly.
- A / B tests show that we cannot reject null hypotheses. In other words, there are no statistically significant differences in conversions between two groups across all 4 events.

External factors:

- The A / B test was carried out on the eve of Christmas during the period of increased buying activity before the Christmas holidays. Thus, we cannot trust the results 100%.

Test result:

- The test group B did not achieve the expected result of increasing total conversion at least 10% on each stage of the funnel product_page → product_cart → purchase.
- Moreover, total conversion in group B is almost the same or less than total conversion in group A on stages product_page -> product_cart.
- Only at the last stage,purchase, test group B showed a 1.11% better result than in group A and than at the product_cart stage in their group. It means that some EU users made purchasing in one-click using the 'Buy Now' button.

Recommendation:

- It can be assumed that EU users prefer to shop in the shortest and fastest way, skipping stages in the usual sales funnel order. This hypothesis should be tested.

Start coding or [generate](#) with AI.

**CHECKOUT THE PORTFOLIO
FOR
COMPLETE WORKBOOK**



THENNARASU M